

# GradHpO

Eynullayev Altay, Rubstsov Denis, Karpeev Gleb

February 26, 2026

## Abstract

This document presents the architecture and design of `gradhpo`, a JAX-based library for short-horizon gradient-based hyperparameter optimization. The library implements four state-of-the-art algorithms: T1-T2 with DARTS approximation, Greedy gradient-based optimization and Online hyperparameter meta-learning with hypergradient distillation. The design follows a state-based approach with a unified API for all algorithms.

## Contents

<b>1</b>	<b>Architecture</b>	<b>2</b>
1.1	Directory Structure . . . . .	2
<b>2</b>	<b>Core Components</b>	<b>2</b>
2.1	Type Definitions . . . . .	2
2.1.1	Core Types ( <code>gradhpo/core/types.py</code> ) . . . . .	2
2.2	Bilevel State . . . . .	3
2.2.1	<code>BilevelState</code> Class ( <code>gradhpo/core/state.py</code> ) . . . . .	3
2.3	Base Optimizer . . . . .	4
2.3.1	<code>BilevelOptimizer</code> Class ( <code>gradhpo/core/base.py</code> ) . . . . .	4
<b>3</b>	<b>Algorithm Implementations</b>	<b>5</b>
3.1	T1-T2 with DARTS . . . . .	5
3.1.1	<code>T1T2Optimizer</code> Class ( <code>gradhpo/algorithms/t1_t2.py</code> ) . . . . .	5
3.2	Greedy Gradient-Based . . . . .	6
3.2.1	<code>GreedyOptimizer</code> Class ( <code>gradhpo/algorithms/greedy.py</code> ) . . . . .	6
3.3	Online with Hypergradient Distillation . . . . .	7
3.3.1	<code>OnlineHypergradientOptimizer</code> Class ( <code>gradhpo/algorithms/online.py</code> ) . . . . .	7
<b>4</b>	<b>Libraries and Dependencies</b>	<b>8</b>
4.1	Integration . . . . .	8

# 1 Architecture

## 1.1 Directory Structure

Path	Description
gradhpo/src	
gradhpo/	Main package directory
core/	Core components
state.py	BilevelState class
base.py	Base optimizer class
types.py	Type definitions
algorithms/	Algorithm implementations
t1_t2.py	T1-T2 with DARTS
greedy.py	Greedy gradient-based
online.py	Online hypergradient distillation
implicit_diff.py	Implicit differentiation
utils/	Utility functions
gradients.py	Gradient utilities
validation.py	Input validation
logging.py	Logging utilities
examples/	Example scripts
tests/	Test suite
docs/	Documentation
setup.py	Package setup script
requirements.txt	Dependencies
README.md	Project readme

Table 1: Project directory structure

## 2 Core Components

### 2.1 Type Definitions

#### 2.1.1 Core Types (gradhpo/core/types.py)

```
1 from typing import Any, Callable, Dict, NamedTuple, Protocol
2 import jax.numpy as jnp
3 from jax import Array
4
5 # Type aliases
6 PyTree = Any # JAX PyTree (nested structure of arrays)
7 LossFn = Callable[[PyTree, PyTree, Any], Array]
8 MetricDict = Dict[str, Union[float, Array]]
9
10 class DataBatch(NamedTuple):
11     """Structure for a data batch.
12
13     Attributes:
14         inputs: Input features [batch_size, ...]
15         targets: Target labels [batch_size, ...]
16     """
17     inputs: Array
18     targets: Array
19
20 class LossFunctions(NamedTuple):
21     """Container for loss functions.
22
23     Attributes:
24         train_loss: Training loss function
```

```

25     val_loss: Validation loss function
26     """
27     train_loss: LossFn
28     val_loss: LossFn

```

## 2.2 Bilevel State

### 2.2.1 BilevelState Class (gradhpo/core/state.py)

The BilevelState class maintains all necessary state information for bilevel optimization.

```

1 @dataclass
2 class BilevelState:
3     """State container for bilevel optimization process.
4
5     Attributes:
6         params: Model parameters (inner level)
7         hyperparams: Hyperparameters to optimize (outer level)
8         inner_opt_state: State of inner optimizer
9         outer_opt_state: State of outer optimizer
10        step: Current optimization step
11        metadata: Additional information (losses, norms, etc.)
12    """
13    params: PyTree
14    hyperparams: PyTree
15    inner_opt_state: OptState
16    outer_opt_state: OptState
17    step: int
18    metadata: Dict[str, Any]

```

#### Key Methods:

```

1 @classmethod
2 def create(cls,
3             params: PyTree,
4             hyperparams: PyTree,
5             inner_opt_state: OptState,
6             outer_opt_state: OptState) -> 'BilevelState':
7     """Create a new BilevelState with initial values.
8
9     Args:
10        params: Initial model parameters
11        hyperparams: Initial hyperparameters
12        inner_opt_state: Initial inner optimizer state
13        outer_opt_state: Initial outer optimizer state
14
15     Returns:
16        Initialized state object
17    """
18
19 def update(self,
20            params: Optional[PyTree] = None,
21            hyperparams: Optional[PyTree] = None,
22            inner_opt_state: Optional[OptState] = None,
23            outer_opt_state: Optional[OptState] = None,
24            metadata: Optional[Dict[str, Any]] = None) -> 'BilevelState':
25     """Create updated state with new values.
26
27     Args:
28        params: New parameters (if None, keep current)
29        hyperparams: New hyperparameters (if None, keep current)
30        inner_opt_state: New inner optimizer state
31        outer_opt_state: New outer optimizer state
32        metadata: New metadata to merge

```

```

33
34     Returns:
35         New state object with updates
36     """
37
38 def get_metric(self, key: str, default: Any = None) -> Any:
39     """Retrieve a metric from metadata.
40
41     Args:
42         key: Metric name
43         default: Default value if key not found
44
45     Returns:
46         Metric value or default
47     """

```

## 2.3 Base Optimizer

### 2.3.1 BilevelOptimizer Class (gradhpo/core/base.py)

Abstract base class for all bilevel optimization algorithms.

```

1 class BilevelOptimizer(ABC):
2     """Abstract base class for bilevel hyperparameter optimizers.
3
4     All algorithms inherit from this class and implement
5     the required abstract methods.
6
7     Attributes:
8         inner_optimizer: Optax optimizer for parameters
9         outer_optimizer: Optax optimizer for hyperparameters
10    """
11
12    def __init__(self,
13                 inner_optimizer: optax.GradientTransformation,
14                 outer_optimizer: optax.GradientTransformation):
15        """Initialize bilevel optimizer."""

```

#### Abstract Methods:

```

1 @abstractmethod
2 def init(self,
3          params: PyTree,
4          hyperparams: PyTree) -> BilevelState:
5     """Initialize the bilevel optimization state.
6
7     Args:
8         params: Initial model parameters
9         hyperparams: Initial hyperparameters
10
11    Returns:
12        Initial optimization state
13    """
14
15 @abstractmethod
16 def step(self,
17          state: BilevelState,
18          train_batch: DataBatch,
19          val_batch: DataBatch,
20          train_loss_fn: LossFn,
21          val_loss_fn: LossFn) -> BilevelState:
22     """Perform one bilevel optimization step.
23
24     Args:

```

```

25     state: Current bilevel state
26     train_batch: Training data batch
27     val_batch: Validation data batch
28     train_loss_fn: Training loss function
29     val_loss_fn: Validation loss function
30
31     Returns:
32         Updated state
33     """
34
35     @abstractmethod
36     def compute_hypergradient(self,
37             state: BilevelState,
38             train_batch: DataBatch,
39             val_batch: DataBatch,
40             train_loss_fn: LossFn,
41             val_loss_fn: LossFn) -> PyTree:
42         """Compute hypergradient w.r.t. hyperparameters.
43
44         Args:
45             state: Current state
46             train_batch: Training batch
47             val_batch: Validation batch
48             train_loss_fn: Training loss
49             val_loss_fn: Validation loss
50
51         Returns:
52             Hypergradient with same structure as hyperparams
53     """

```

## 3 Algorithm Implementations

### 3.1 T1-T2 with DARTS

#### 3.1.1 T1T2Optimizer Class (gradhpo/algorithms/t1\_t2.py)

Implements T1-T2 algorithm using DARTS-style finite difference approximation.

```

1  class T1T2Optimizer(BilevelOptimizer):
2      """T1-T2 algorithm with DARTS approximation.
3
4      Attributes:
5          inner_optimizer: Optimizer for model parameters
6          outer_optimizer: Optimizer for hyperparameters
7          inner_steps: Number of inner loop steps (K)
8          epsilon: Step size for finite difference approximation
9      """
10
11     def __init__(self,
12             inner_optimizer: optax.GradientTransformation,
13             outer_optimizer: optax.GradientTransformation,
14             inner_steps: int = 1,
15             epsilon: float = 0.01):
16         """Initialize T1-T2 optimizer.
17
18         Args:
19             inner_optimizer: Optax optimizer for parameters
20             outer_optimizer: Optax optimizer for hyperparameters
21             inner_steps: Number of inner optimization steps
22             epsilon: Finite difference epsilon
23     """

```

**Key Methods:**

```

1 def init(self,
2         params: PyTree,
3         hyperparams: PyTree) -> BilevelState:
4     """Initialize T1-T2 optimization state."""
5
6 @jit
7 def step(self,
8          state: BilevelState,
9          train_batch: DataBatch,
10         val_batch: DataBatch,
11         train_loss_fn: LossFn,
12         val_loss_fn: LossFn) -> BilevelState:
13     """Perform one T1-T2 optimization step.
14
15     Algorithm:
16         1. Update parameters on training loss (inner loop)
17         2. Compute hypergradient using DARTS approximation
18         3. Update hyperparameters
19     """
20
21 def compute_hypergradient(self,
22                          state: BilevelState,
23                          train_batch: DataBatch,
24                          val_batch: DataBatch,
25                          train_loss_fn: LossFn,
26                          val_loss_fn: LossFn) -> PyTree:
27     """Compute hypergradient using finite difference.
28
29     Uses DARTS-style approximation from Section 2.3.
30     """

```

## 3.2 Greedy Gradient-Based

### 3.2.1 GreedyOptimizer Class (gradhpo/algorithms/greedy.py)

Implements generalized greedy gradient-based hyperparameter optimization.

```

1 class GreedyOptimizer(BilevelOptimizer):
2     """Greedy gradient-based hyperparameter optimization.
3
4     Attributes:
5         inner_optimizer: Optimizer for parameters
6         outer_optimizer: Optimizer for hyperparameters
7         unroll_steps: Number of steps to unroll
8     """
9
10    def __init__(self,
11                 inner_optimizer: optax.GradientTransformation,
12                 outer_optimizer: optax.GradientTransformation,
13                 unroll_steps: int = 1):
14        """Initialize Greedy optimizer.
15
16        Args:
17            inner_optimizer: Optimizer for parameters
18            outer_optimizer: Optimizer for hyperparameters
19            unroll_steps: Number of inner steps to unroll
20        """

```

#### Key Methods:

```

1 def init(self,
2         params: PyTree,
3         hyperparams: PyTree) -> BilevelState:
4     """Initialize Greedy optimization state."""

```

```

5
6 @jit
7 def step(self,
8         state: BilevelState,
9         train_batch: DataBatch,
10        val_batch: DataBatch,
11        train_loss_fn: LossFn,
12        val_loss_fn: LossFn) -> BilevelState:
13     """Perform one greedy optimization step."""
14
15 def compute_hypergradient(self,
16                          state: BilevelState,
17                          train_batch: DataBatch,
18                          val_batch: DataBatch,
19                          train_loss_fn: LossFn,
20                          val_loss_fn: LossFn) -> PyTree:
21     """Compute hypergradient by unrolling inner optimization."""

```

### 3.3 Online with Hypergradient Distillation

#### 3.3.1 OnlineHypergradientOptimizer Class (gradhpo/algorithms/online.py)

Implements online hyperparameter meta-learning with hypergradient distillation.

```

1 class OnlineHypergradientOptimizer(BilevelOptimizer):
2     """Online optimization with hypergradient distillation.
3
4     Attributes:
5         inner_optimizer: Optimizer for parameters
6         outer_optimizer: Optimizer for hyperparameters
7         distillation_weight: Weight for distillation loss
8         memory_size: Number of past hypergradients to store
9     """
10
11    def __init__(self,
12                 inner_optimizer: optax.GradientTransformation,
13                 outer_optimizer: optax.GradientTransformation,
14                 distillation_weight: float = 0.1,
15                 memory_size: int = 10):
16        """Initialize Online Hypergradient optimizer.
17
18        Args:
19            inner_optimizer: Optimizer for parameters
20            outer_optimizer: Optimizer for hyperparameters
21            distillation_weight: Weight for distillation term
22            memory_size: Size of hypergradient memory buffer
23    """

```

#### Key Methods:

```

1 def init(self,
2          params: PyTree,
3          hyperparams: PyTree) -> BilevelState:
4     """Initialize state with hypergradient memory."""
5
6 @jit
7 def step(self,
8          state: BilevelState,
9          train_batch: DataBatch,
10         val_batch: DataBatch,
11         train_loss_fn: LossFn,
12         val_loss_fn: LossFn) -> BilevelState:
13     """Perform one online optimization step with distillation."""
14

```

```

15 def compute_hypergradient(self,
16     state: BilevelState,
17     train_batch: DataBatch,
18     val_batch: DataBatch,
19     train_loss_fn: LossFn,
20     val_loss_fn: LossFn) -> PyTree:
21     """Compute hypergradient using short-horizon approximation."""

```

## 4 Libraries and Dependencies

Library	Version
JAX	$\geq 0.4.0$
jaxlib	$\geq 0.4.0$
optax	$\geq 0.1.7$
chex	$\geq 0.1.8$
typing-extensions	$\geq 4.5.0$
numpy	$\geq 1.24.0$

Table 2: Core dependencies

### 4.1 Integration

- **JAX**: Core framework for automatic differentiation and JIT compilation
- **Optax**: Provides gradient transformation and optimization algorithms
- **Chex**: Testing utilities for JAX code (assertions, variants)
- **NumPy**: Array operations and compatibility layer